

2203A51530 ML ASSIGNMENT

```
import pandas as pd
import seaborn as sns
import os
import numpy as np
import matplotlib.pyplot as plt

housing_df = pd.read_csv('/content/housing.csv')

# Use .info() to show the features (i.e. columns) in your dataset
along with a count and datatype
housing_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

housing_df.shape

(20640, 10)

housing_df.head()

{"summary":{"\n  \"name\": \"housing_df\",\n  \"rows\": 20640,\n  \"fields\": [\n    {\n      \"column\": \"longitude\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2.0035317235025882,\n        \"min\": -124.35,\n        \"max\": -114.31,\n        \"num_unique_values\": 844,\n        \"samples\": [\n          -118.63,\n          -119.86,\n          -121.26\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"latitude\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2.1359523974571153,\n        \"min\": 32.54,\n        \"max\": 41.95,\n        \"num_unique_values\": 862,\n        \"samples\": [\n
```

```

33.7,\n          34.41,\n          38.24\n          ],\n  \"semantic_type\": \"\",\n  \"description\": \"\"\n  },\n  {\n    \"column\": \"housing_median_age\",\n    \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 12.58555761211165,\n      \"min\": 1.0,\n      \"max\": 52.0,\n      \"num_unique_values\": 52,\n      \"samples\": [\n        35.0,\n        25.0,\n        7.0\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\"\n    },\n    {\n      \"column\": \"total_rooms\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2181.615251582795,\n        \"min\": 2.0,\n        \"max\": 39320.0,\n        \"num_unique_values\": 5926,\n        \"samples\": [\n          699.0,\n          1544.0,\n          3966.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\": \"total_bedrooms\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 421.3850700740323,\n          \"min\": 1.0,\n          \"max\": 6445.0,\n          \"num_unique_values\": 1923,\n          \"samples\": [\n            1538.0,\n            1298.0,\n            1578.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\": \"population\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 1132.462121765341,\n            \"min\": 3.0,\n            \"max\": 35682.0,\n            \"num_unique_values\": 3888,\n            \"samples\": [\n              4169.0,\n              636.0,\n              3367.0\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          },\n          {\n            \"column\": \"households\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 382.32975283161073,\n              \"min\": 1.0,\n              \"max\": 6082.0,\n              \"num_unique_values\": 1815,\n              \"samples\": [\n                21.0,\n                750.0,\n                1447.0\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            },\n            {\n              \"column\": \"median_income\",\n              \"properties\": {\n                \"dtype\": \"number\",\n                \"std\": 1.8998217179452688,\n                \"min\": 0.4999,\n                \"max\": 15.0001,\n                \"num_unique_values\": 12928,\n                \"samples\": [\n                  5.0286,\n                  2.0433,\n                  6.1228\n                ],\n                \"semantic_type\": \"\",\n                \"description\": \"\"\n              },\n              {\n                \"column\": \"median_house_value\",\n                \"properties\": {\n                  \"dtype\": \"number\",\n                  \"std\": 115395.61587441387,\n                  \"min\": 14999.0,\n                  \"max\": 500001.0,\n                  \"num_unique_values\": 3842,\n                  \"samples\": [\n                    194300.0,\n                    379000.0,\n                    230100.0\n                  ],\n                  \"semantic_type\": \"\",\n                  \"description\": \"\"\n                },\n                {\n                  \"column\": \"ocean_proximity\",\n                  \"properties\": {\n                    \"dtype\": \"category\",\n                    \"num_unique_values\": 5,\n                    \"samples\": [\n                      \"<1H OCEAN\",\n                      \"ISLAND\",\n                      \"INLAND\"\n                    ],\n                    \"semantic_type\": \"\",\n                    \"description\": \"\"\n                  }\n                }\n              ],\n            \"type\": \"dataframe\", \"variable_name\": \"housing_df\"

```

```
housing_df.tail()
```

```
{
  "summary": {
    "name": "housing_df",
    "rows": 5,
    "fields": [
      {
        "column": "longitude",
        "properties": {
          "dtype": "number",
          "std": 0.08264381404557382,
          "min": -121.32,
          "max": -121.09,
          "num_unique_values": 5,
          "samples": [
            -121.21,
            -121.24,
            -121.22
          ]
        },
        "semantic_type": "number",
        "description": "Longitude"
      },
      {
        "column": "latitude",
        "properties": {
          "dtype": "number",
          "std": 0.047958315233128025,
          "min": 39.37,
          "max": 39.49,
          "num_unique_values": 4,
          "samples": [
            39.49,
            39.37,
            39.48
          ]
        },
        "semantic_type": "number",
        "description": "Latitude"
      },
      {
        "column": "housing_median_age",
        "properties": {
          "dtype": "number",
          "std": 3.5637059362410923,
          "min": 16.0,
          "max": 25.0,
          "num_unique_values": 4,
          "samples": [
            16.0,
            25.0
          ]
        },
        "semantic_type": "number",
        "description": "Housing median age"
      },
      {
        "column": "total_rooms",
        "properties": {
          "dtype": "number",
          "std": 774.7823565363373,
          "min": 697.0,
          "max": 2785.0,
          "num_unique_values": 5,
          "samples": [
            697.0,
            2785.0,
            2254.0
          ]
        },
        "semantic_type": "number",
        "description": "Total rooms"
      },
      {
        "column": "total_bedrooms",
        "properties": {
          "dtype": "number",
          "std": 170.95818202121828,
          "min": 150.0,
          "max": 616.0,
          "num_unique_values": 5,
          "samples": [
            150.0,
            616.0,
            485.0
          ]
        },
        "semantic_type": "number",
        "description": "Total bedrooms"
      },
      {
        "column": "population",
        "properties": {
          "dtype": "number",
          "std": 376.6566075352987,
          "min": 356.0,
          "max": 1387.0,
          "num_unique_values": 5,
          "samples": [
            356.0,
            1387.0,
            1007.0
          ]
        },
        "semantic_type": "number",
        "description": "Population"
      },
      {
        "column": "households",
        "properties": {
          "dtype": "number",
          "std": 154.41729177783168,
          "min": 114.0,
          "max": 530.0,
          "num_unique_values": 5,
          "samples": [
            114.0,
            530.0,
            433.0
          ]
        },
        "semantic_type": "number",
        "description": "Households"
      },
      {
        "column": "median_income",
        "properties": {
          "dtype": "number",
          "std": 0.43616087857578417,
          "min": 1.5603,
          "max": 2.5568,
          "num_unique_values": 5,
          "samples": [
            2.5568,
            2.3886,
            1.7
          ]
        },
        "semantic_type": "number",
        "description": "Median income"
      }
    ]
  }
}
```

```

n    },\n    {\n        \"column\": \"median_house_value\", \n        \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 6716.546731766258, \n            \"min\": 77100.0, \n            \"max\": 92300.0, \n            \"num_unique_values\": 5, \n            \"samples\": [\n                77100.0, \n                89400.0, \n                92300.0\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        } \n    }, \n    {\n        \"column\": \"ocean_proximity\", \n        \"properties\": {\n            \"dtype\": \"category\", \n            \"num_unique_values\": 1, \n            \"samples\": [\n                \"INLAND\" \n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        } \n    } \n    ], \n    \"type\": \"dataframe\"}

```

```
housing_df.describe()
```

```

{\"summary\": \"{ \n    \"name\": \"housing_df\", \n    \"rows\": 8, \n    \"fields\": [ \n        { \n            \"column\": \"longitude\", \n            \"properties\": { \n                \"dtype\": \"number\", \n                \"std\": 7333.554670164394, \n                \"min\": -124.35, \n                \"max\": 20640.0, \n                \"num_unique_values\": 8, \n                \"samples\": [ \n                    -119.56970445736432, \n                    -118.49, \n                    20640.0 \n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\" \n            } \n        }, \n        { \n            \"column\": \"latitude\", \n            \"properties\": { \n                \"dtype\": \"number\", \n                \"std\": 2.1359523974571153, \n                \"min\": 35.63186143410853, \n                \"max\": 34.26, \n                \"num_unique_values\": 8, \n                \"samples\": [ \n                    20640.0 \n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\" \n            } \n        }, \n        { \n            \"column\": \"housing_median_age\", \n            \"properties\": { \n                \"dtype\": \"number\", \n                \"std\": 7288.35672120143, \n                \"min\": 1.0, \n                \"max\": 20640.0, \n                \"num_unique_values\": 8, \n                \"samples\": [ \n                    28.639486434108527, \n                    29.0, \n                    20640.0 \n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\" \n            } \n        }, \n        { \n            \"column\": \"total_rooms\", \n            \"properties\": { \n                \"dtype\": \"number\", \n                \"std\": 13944.990983306392, \n                \"min\": 2.0, \n                \"max\": 39320.0, \n                \"num_unique_values\": 8, \n                \"samples\": [ \n                    2635.7630813953488, \n                    2127.0, \n                    20640.0 \n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\" \n            } \n        }, \n        { \n            \"column\": \"total_bedrooms\", \n            \"properties\": { \n                \"dtype\": \"number\", \n                \"std\": 7106.427031043753, \n                \"min\": 1.0, \n                \"max\": 20433.0, \n                \"num_unique_values\": 8, \n                \"samples\": [ \n                    537.8705525375618, \n                    435.0, \n                    20433.0 \n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\" \n            } \n        }, \n        { \n            \"column\": \"population\", \n            \"properties\": { \n                \"dtype\": \"number\", \n                \"std\": 13192.258841737372, \n                \"min\": 3.0, \n                \"max\": 35682.0, \n                \"num_unique_values\": 8, \n                \"samples\": [ \n

```

```

1425.4767441860465,\n          1166.0,\n          20640.0\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\"\n  },\n  {\n    \"column\": \"households\",\n    \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 7167.532601135343,\n      \"min\": 1.0,\n      \"max\": 20640.0,\n      \"num_unique_values\": 8,\n      \"samples\": [\n        499.5396802325581,\n        409.0,\n        20640.0\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\"\n    },\n    {\n      \"column\": \"median_income\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 7295.7214358536385,\n        \"min\": 0.4999,\n        \"max\": 20640.0,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          3.8706710029069766,\n          3.5347999999999997,\n          20640.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\": \"median_house_value\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 14999.0,\n          \"min\": 156160.28379826446,\n          \"max\": 500001.0,\n          \"num_unique_values\": 8,\n          \"samples\": [\n            179700.0,\n            20640.0,\n            206855.81690891474,\n            179700.0,\n            20640.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\": \"ocean_proximity\",\n          \"properties\": {\n            \"dtype\": \"object\",\n            \"std\": 0,\n            \"min\": 0,\n            \"max\": 0,\n            \"num_unique_values\": 1,\n            \"samples\": [\n              1.002906976744186\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          }\n        }\n      }\n    },\n    {\n      \"column\": \"total_rooms\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1166.0,\n        \"min\": 1,\n        \"max\": 20640.0,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          1425.4767441860465,\n          1166.0,\n          20640.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  }\n],\n\"type\": \"dataframe\"}

```

```
housing_df.isnull().sum()
```

```

longitude      0
latitude       0
housing_median_age  0
total_rooms    0
total_bedrooms 207
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity 0
dtype: int64

```

```
# Calculate the % of missing data
```

```

housing_df['total_bedrooms'].isnull().sum()/housing_df.shape[0] * 100
1.002906976744186

```

```
from sklearn.impute import KNNImputer
```

```
# create a temporary copy of the dataset
```

```
housing_df_temp = housing_df.copy()
```

```
# retrieve columns with numerical data; will exclude the
ocean_proximity column since the datatype is object; other columns are
float64
```

```
columns_list = [col for col in housing_df_temp.columns if
```

```

housing_df_temp[col].dtype != 'object']

# extract columns that contain at least one missing value
new_column_list = [col for col in housing_df_temp.loc[:,
housing_df_temp.isnull().any()]]

# update temp dataframe with numeric columns that have empty values
housing_df_temp = housing_df_temp[new_column_list]

# initialize KNNImputer to impute missing data using machine learning
knn = KNNImputer(n_neighbors = 3)

# fit function trains the model
knn.fit(housing_df_temp)

# transform the data using the model
# applies the transformation model (ie knn) to data
array_Values = knn.transform(housing_df_temp)

# convert the array values to a dataframe with the appropriate column
names
housing_df_temp = pd.DataFrame(array_Values, columns =
new_column_list)

# confirm there are no columns with missing values
housing_df_temp.isnull().sum()

total_bedrooms      0
dtype: int64

# overlay the imputed column over the old column with missing values

# loop through the list of columns and overlay each one
for column_name in new_column_list:
    housing_df[column_name] =
housing_df_temp.replace(housing_df[column_name],housing_df[column_name
])

# confirm columns no longer contain null data
housing_df.isnull().sum()

longitude           0
latitude            0
housing_median_age  0
total_rooms         0
total_bedrooms      0
population          0
households          0
median_income       0
median_house_value  0

```

```

ocean_proximity      0
dtype: int64


# Plot the distribution of the target variable (median_house_value)
using a histogram

# bins->amount of columns
plt.hist(housing_df['median_house_value'], bins=80)
plt.xlabel("House Values")

# We can see from the plot that the values of Median House Value are
distributed normally with few outliers.
# Most of the house are around 100,000-200,000 range

Text(0.5, 0, 'House Values')

```

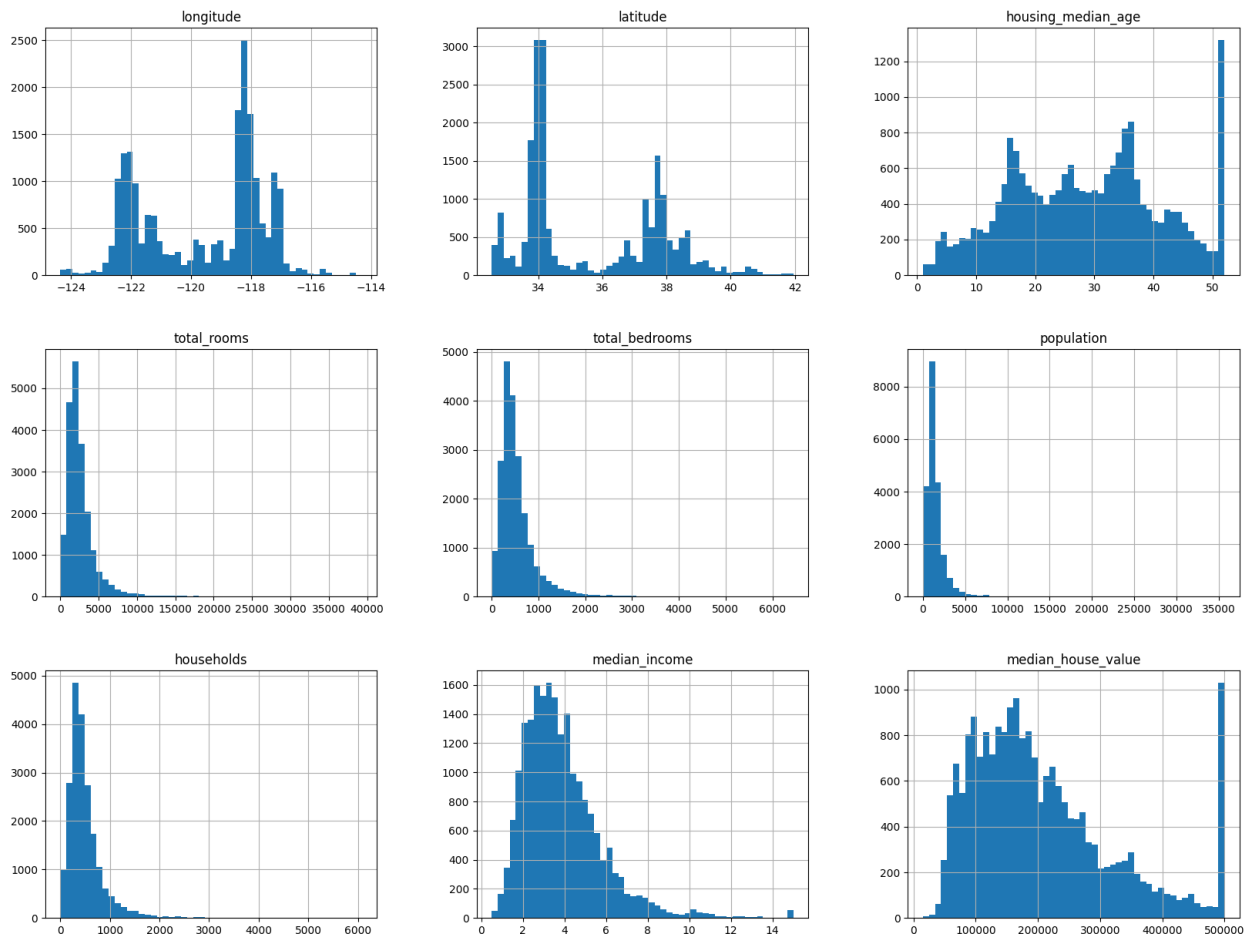


```

# let's do histograms for the all the features to understand the data
distributions
# using housing_df as to not plot the encoded values for
OCEAN_PROXIMITY
housing_df.hist(bins=50, figsize=(20,15))

```

```
array([[<Axes: title={'center': 'longitude'}>,
       <Axes: title={'center': 'latitude'}>,
       <Axes: title={'center': 'housing_median_age'}>],
      [<Axes: title={'center': 'total_rooms'}>,
       <Axes: title={'center': 'total_bedrooms'}>,
       <Axes: title={'center': 'population'}>],
      [<Axes: title={'center': 'households'}>,
       <Axes: title={'center': 'median_income'}>,
       <Axes: title={'center': 'median_house_value'}>]],
dtype=object)
```



```
# Plot a graphical correlation matrix for each pair of columns in the
dataframe
```

```
corr = housing_df.corr() # data frame correlation function
print(corr)
```

	longitude	latitude	housing_median_age
total_rooms \			
longitude	1.000000	-0.924664	-0.108197
0.044568			
latitude	-0.924664	1.000000	0.011173
			-


```

0.036100
housing_median_age  -0.108197  0.011173          1.000000  -
0.361262
total_rooms          0.044568 -0.036100          -0.361262
1.000000
total_bedrooms       0.069260 -0.066658          -0.318998
0.927253
population           0.099773 -0.108785          -0.296244
0.857126
households           0.055310 -0.071035          -0.302916
0.918484
median_income        -0.015176 -0.079809          -0.119034
0.198050
median_house_value   -0.045967 -0.144160          0.105623
0.134153

```

```

                                total_bedrooms  population  households
median_income \
longitude          0.069260    0.099773    0.055310    -
0.015176
latitude          -0.066658    -0.108785    -0.071035    -
0.079809
housing_median_age -0.318998    -0.296244    -0.302916    -
0.119034
total_rooms        0.927253    0.857126    0.918484
0.198050
total_bedrooms     1.000000    0.873910    0.974725    -
0.007682
population         0.873910    1.000000    0.907222
0.004834
households         0.974725    0.907222    1.000000
0.013033
median_income      -0.007682    0.004834    0.013033
1.000000
median_house_value  0.049454    -0.024650    0.065843
0.688075

```

```

                                median_house_value
longitude          -0.045967
latitude          -0.144160
housing_median_age  0.105623
total_rooms        0.134153
total_bedrooms     0.049454
population         -0.024650
households         0.065843
median_income      0.688075
median_house_value 1.000000

```

```

<ipython-input-15-3abd71ce2464>:2: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it

```

will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.


```
corr = housing_df.corr() # data frame correlation function
```

make the heatmap larger in size

```
plt.figure(figsize = (8,8))
```

```
sns.heatmap(corr, annot=True)
```

```
plt.show()
```



```

# Additionally we noted that several features
(total_rooms, total_bedrooms, population, households) have very high
correlation to one another,
# so it's interesting to find out if a removal of a few of them would
have any affect on the model performance

# a new feature that is a ratio of the total rooms to households
housing_df['rooms_per_household'] =
housing_df['total_rooms']/housing_df['households']

# a new feature that is a ratio of the total bedrooms to the total
rooms
housing_df['bedrooms_per_room'] =
housing_df['total_bedrooms']/housing_df['total_rooms']

# a new feature that is a ratio of the population to the households
housing_df['population_per_household']=
housing_df['population']/housing_df['households']

# let's combine the latitude and longitude into 1
housing_df['coords'] = housing_df['longitude']/housing_df['latitude']

housing_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   longitude                             20640 non-null  float64
1   latitude                             20640 non-null  float64
2   housing_median_age                    20640 non-null  float64
3   total_rooms                           20640 non-null  float64
4   total_bedrooms                       20640 non-null  float64
5   population                            20640 non-null  float64
6   households                            20640 non-null  float64
7   median_income                        20640 non-null  float64
8   median_house_value                   20640 non-null  float64
9   ocean_proximity                      20640 non-null  object
10  rooms_per_household                   20640 non-null  float64
11  bedrooms_per_room                     20640 non-null  float64
12  population_per_household              20640 non-null  float64
13  coords                               20640 non-null  float64
dtypes: float64(13), object(1)
memory usage: 2.2+ MB

# remove total_rooms, households, total bedrooms, popluation,
longitude, latitude
housing_df = housing_df.drop('total_rooms', axis=1)
housing_df = housing_df.drop('households', axis=1)

```

```
housing_df = housing_df.drop('total_bedrooms', axis=1)
housing_df = housing_df.drop('population', axis=1)
housing_df = housing_df.drop('longitude', axis=1)
housing_df = housing_df.drop('latitude', axis=1)
```

```
housing_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20640 entries, 0 to 20639
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	housing_median_age	20640 non-null	float64
1	median_income	20640 non-null	float64
2	median_house_value	20640 non-null	float64
3	ocean_proximity	20640 non-null	object
4	rooms_per_household	20640 non-null	float64
5	bedrooms_per_room	20640 non-null	float64
6	population_per_household	20640 non-null	float64
7	coords	20640 non-null	float64

```
dtypes: float64(7), object(1)
```

```
memory usage: 1.3+ MB
```

```
#Heatmap after removing correlation
```

```
corr = housing_df.corr()
```

```
#make the heatmap larger in size
```


```
plt.figure(figsize = (7,7))
```

```
sns.heatmap(corr, annot=True)
```

```
plt.show()
```

```
<ipython-input-19-1264607259b1>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
```

```
corr = housing_df.corr()
```



```
#Encoding categorical data
# Most ML algorithms can only learn from numeric data (it's all Math)
so categorical data must be encoded (i.e. converted) to numeric data

# Let's review our data types again; showing that ocean_proximity is
the only categorical data
housing_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
#   ...
#   ocean_proximity        20640 non-null  object
```

```

---
0  housing_median_age    20640 non-null float64
1  median_income         20640 non-null float64
2  median_house_value    20640 non-null float64
3  ocean_proximity       20640 non-null object
4  rooms_per_household   20640 non-null float64
5  bedrooms_per_room     20640 non-null float64
6  population_per_household 20640 non-null float64
7  coords                20640 non-null float64

```

```

dtypes: float64(7), object(1)
memory usage: 1.3+ MB

```

```

# let's see the unique categories for OCEAN_PROXIMITY

```

```

housing_df.ocean_proximity.unique()

```

```

array(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND'],
      dtype=object)

```

```

# let's count

```

```

housing_df["ocean_proximity"].value_counts()

```

```

<1H OCEAN    9136
INLAND       6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND         5

```

```

Name: ocean_proximity, dtype: int64

```

```

# Let's see how the Panda's get_dummies() function works (generates
new columns based on the possible options)

```

```

print(pd.get_dummies(housing_df['ocean_proximity']))

```

	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
0	0	0	0	1	0
1	0	0	0	1	0
2	0	0	0	1	0
3	0	0	0	1	0
4	0	0	0	1	0
...
20635	0	1	0	0	0
20636	0	1	0	0	0
20637	0	1	0	0	0
20638	0	1	0	0	0
20639	0	1	0	0	0

```

[20640 rows x 5 columns]

```

```

# let's replace the OCEAN_PROXIMITY column using get_dummies()

```

```

housing_df_encoded = pd.get_dummies(data=housing_df,
columns=['ocean_proximity'])

```

```
# print the first few observations; notice the old OCEAN_PROXIMITY
column is gone
```

```
housing_df_encoded.head()
```

```
{"summary":{"\n  \"name\": \"housing_df_encoded\",\n  \"rows\": 20640,\n  \"fields\": [\n    {\n      \"column\": \"housing_median_age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 12.58555761211165,\n        \"min\": 1.0,\n        \"max\": 52.0,\n        \"num_unique_values\": 52,\n        \"samples\": [\n          35.0,\n          25.0,\n          7.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"median_income\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1.8998217179452688,\n        \"min\": 0.4999,\n        \"max\": 15.0001,\n        \"num_unique_values\": 12928,\n        \"samples\": [\n          5.0286,\n          2.0433,\n          6.1228\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"median_house_value\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 115395.61587441387,\n        \"min\": 14999.0,\n        \"max\": 500001.0,\n        \"num_unique_values\": 3842,\n        \"samples\": [\n          194300.0,\n          379000.0,\n          230100.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"rooms_per_household\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2.4741731394243187,\n        \"min\": 0.8461538461538461,\n        \"max\": 141.9090909090909,\n        \"num_unique_values\": 19392,\n        \"samples\": [\n          6.111269614835948,\n          5.912820512820513,\n          5.7924528301886795\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"bedrooms_per_room\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.06976426445426125,\n        \"min\": 0.045936506323132786,\n        \"max\": 3.4926659255685832,\n        \"num_unique_values\": 19473,\n        \"samples\": [\n          0.19223760932944606,\n          0.2103494623655914,\n          0.20042417815482502\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"population_per_household\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 10.386049562213618,\n        \"min\": 0.6923076923076923,\n        \"max\": 1243.3333333333333,\n        \"num_unique_values\": 18841,\n        \"samples\": [\n          2.6939799331103678,\n          3.559375,\n          3.297082228116711\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"coords\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.1468967022257624,\n        \"min\": -3.597235023041475,\n        \"max\": -2.87341469251017,\n        \"num_unique_values\": 12575,\n        \"samples\": [\n          -3.1434118560704114,\n          -3.297632158590308,\n          -3.475270705297044\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}
```

```

n    },\n    {\n        \"column\": \"ocean_proximity_<1H OCEAN\", \n        \"properties\": {\n            \"dtype\": \"uint8\", \n            \"num_unique_values\": 2, \n            \"samples\": [\n                1, \n                0\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        } \n    }, \n    {\n        \"column\": \"ocean_proximity_INLAND\", \n        \"properties\": {\n            \"dtype\": \"uint8\", \n            \"num_unique_values\": 2, \n            \"samples\": [\n                1, \n                0\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        } \n    }, \n    {\n        \"column\": \"ocean_proximity_ISLAND\", \n        \"properties\": {\n            \"dtype\": \"uint8\", \n            \"num_unique_values\": 2, \n            \"samples\": [\n                1, \n                0\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        } \n    }, \n    {\n        \"column\": \"ocean_proximity_NEAR BAY\", \n        \"properties\": {\n            \"dtype\": \"uint8\", \n            \"num_unique_values\": 2, \n            \"samples\": [\n                0, \n                1\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        } \n    }, \n    {\n        \"column\": \"ocean_proximity_NEAR OCEAN\", \n        \"properties\": {\n            \"dtype\": \"uint8\", \n            \"num_unique_values\": 2, \n            \"samples\": [\n                1, \n                0\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n        } \n    } \n]\n},\n\"type\": \"dataframe\", \"variable_name\": \"housing_df_encoded\"}

```

#Train the model

```
import sklearn
```

```
from sklearn.model_selection import train_test_split
```

remove spaces from column names and convert all to lowercase and remove special characters as it could cause issues in the future

```
housing_df_encoded.columns = [c.lower().replace(' ', '_').replace('<', '_') for c in housing_df_encoded.columns]
```

Split target variable and feature variables

```
X = housing_df_encoded[['housing_median_age',
'median_income', 'bedrooms_per_room', 'population_per_household', 'coords',
'ocean_proximity_1h_ocean',
```

```
'ocean_proximity_inland', 'ocean_proximity_island', 'ocean_proximity_near_bay', 'ocean_proximity_near_ocean']]
```

```
y = housing_df_encoded['median_house_value']
```

```
print(X)
```

	housing_median_age	median_income	bedrooms_per_room	\
0	41.0	8.3252	0.146591	
1	21.0	8.3014	0.155797	
2	52.0	7.2574	0.129516	
3	52.0	5.6431	0.184458	

4	52.0	3.8462	0.172096
...
20635	25.0	1.5603	0.224625
20636	18.0	2.5568	0.215208
20637	17.0	1.7000	0.215173
20638	18.0	1.8672	0.219892
20639	16.0	2.3886	0.221185

population_per_household		coords	
ocean_proximity__lh_ocean \			
0	2.555556	-3.226769	0
1	2.109842	-3.228209	0
2	2.802260	-3.229590	0
3	2.547945	-3.229855	0
4	2.181467	-3.229855	0
...
20635	2.560606	-3.067123	0
20636	3.122807	-3.069385	0
20637	2.325635	-3.074309	0
20638	2.123209	-3.076845	0
20639	2.616981	-3.079502	0

ocean_proximity_inland		ocean_proximity_island \	
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
...	
20635	1	0	
20636	1	0	
20637	1	0	
20638	1	0	
20639	1	0	

ocean_proximity_near_bay		ocean_proximity_near_ocean	
0	1	0	
1	1	0	
2	1	0	
3	1	0	

4	1	0
...
20635	0	0
20636	0	0
20637	0	0
20638	0	0
20639	0	0

[20640 rows x 10 columns]

```
# Split training & test data
# Splitting the data into training and testing sets in numpy arrays
# We train the model with 70% of the samples and test with the
  remaining 30%
# X -> array with the inputs; y -> array of the outputs
X_train, X_test, y_train, y_test = train_test_split(X, y,
  random_state=42, shuffle=True, test_size=0.3)

# Confirm how the data was split
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(14448, 10)
(6192, 10)
(14448,)
(6192,)

#Linear Regression - Model Training
# Use scikit-learn's LinearRegression to train the model on both the
  training and evaluate it on the test sets
from sklearn.linear_model import LinearRegression

# Create a Linear regressor using all the feature variables
reg_model = LinearRegression()

# Train the model using the training sets
reg_model.fit(X_train, y_train)

LinearRegression()

#run the predictions on the training and testing data
y_pred_test = reg_model.predict(X_test)

#compare the actual values (ie, target) with the values predicted by
  the model
pred_test_df = pd.DataFrame({'Actual': y_test, 'Predicted':
  y_pred_test})

pred_test_df
```

```
{
  "summary": {
    "name": "pred_test_df",
    "rows": 6192,
    "fields": [
      {
        "column": "Actual",
        "properties": {
          "dtype": "number",
          "std": 114575.39507173728,
          "min": 14999.0,
          "max": 500001.0,
          "num_unique_values": 2689,
          "samples": [
            203300.0,
            202200.0,
            271100.0
          ],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "Predicted",
        "properties": {
          "dtype": "number",
          "std": 92431.02486873654,
          "min": -7939.178494427237,
          "max": 1461440.4077402034,
          "num_unique_values": 6192,
          "samples": [
            122657.7869513371,
            72978.69492618677,
            207207.3225332344
          ],
          "semantic_type": "",
          "description": ""
        }
      }
    ],
    "type": "dataframe",
    "variable_name": "pred_test_df"
  }
}
```

Determine accuracy using R^2

R^2 : R squared is another way to evaluate the performance of a regression model.

1, means that the model is perfect and 0 means the the model will perform poorly.

```
r2_reg_model_test = round(reg_model.score(X_test, y_test),2)
```

```
print("R^2 Test: {}".format(r2_reg_model_test))
```

```
R^2 Test: 0.56
```

try another machine learning algorithm : Random Forest

Use scikit-learn's Random Forest to train the model on both the training and evaluate it on the test sets

```
from sklearn.ensemble import RandomForestRegressor
```

Create a regressor using all the feature variables

```
rf_model = RandomForestRegressor(n_estimators=10, random_state=10)
```

Train the model using the training sets

```
rf_model.fit(X_train, y_train)
```

```
RandomForestRegressor(n_estimators=10, random_state=10)
```

#run the predictions on the training and testing data

```
y_rf_pred_test = rf_model.predict(X_test)
```

#compare the actual values (ie, target) with the values predicted by the model

```
rf_pred_test_df = pd.DataFrame({'Actual': y_test, 'Predicted': y_rf_pred_test})
```

```
rf_pred_test_df
```

```
{
  "summary": {
    "name": "rf_pred_test_df",
    "rows": 6192,
    "fields": [
      {
        "column": "Actual",
        "properties": {
          "dtype": "number",
          "std": 114575.39507173728,
          "min": 14999.0,
          "max": 500001.0,
          "num_unique_values": 2689,
          "samples": [
            203300.0,
            202200.0,
            271100.0
          ],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "Predicted",
        "properties": {
          "dtype": "number",
          "std": 102994.30885837866,
          "min": 46270.0,
          "max": 500001.0,
          "num_unique_values": 5615,
          "samples": [
            58930.0,
            390480.5,
            85400.0
          ],
          "semantic_type": "",
          "description": ""
        }
      }
    ]
  },
  "type": "dataframe",
  "variable_name": "rf_pred_test_df"
}
```

```
# Determine accuracy using R^2
```

```
from sklearn.metrics import r2_score, mean_squared_error
```

```
score = r2_score(y_test, y_rf_pred_test)
```

```
print("R^2 - {}".format(round(score, 2) * 100))
```

```
R^2 - 75.0%
```

```
# Determine RMSE - Root Mean Squared Error on the test data
```

```
print('RMSE on test data: ', mean_squared_error(y_test,
y_rf_pred_test)**(0.5))
```


```
RMSE on test data: 57289.11495447338
```

```
# Determine feature importance - random forest algorithm is that it
gives you the 'feature importance' for all the variables in the data
# plot the 6 most important features
```

```
plt.figure(figsize=(10,6))
```

```
feat_importances = pd.Series(rf_model.feature_importances_, index =
X_train.columns)
```

```
feat_importances.nlargest(6).plot(kind='barh');
```



```
# training data with 5 most important features
train_x_if = X_train[['bedrooms_per_room', 'housing_median_age',
'coords',
'ocean_proximity_inland', 'population_per_household', 'median_income']]
test_x_if = X_test[['bedrooms_per_room', 'housing_median_age',
'coords',
'ocean_proximity_inland', 'population_per_household', 'median_income']]
```

```
# create an object of the RandomForestRegressor Model
rf_model_if = RandomForestRegressor(n_estimators=10, random_state=10)
```

```
# fit the model with the training data
rf_model_if.fit(train_x_if, y_train)
```

```
# predict the target on the test data
predict_test_with_if = rf_model_if.predict(test_x_if)
```

```
# Root Mean Squared Error on the train and test data
print('RMSE on test data: ', mean_squared_error(y_test,
predict_test_with_if)**(0.5))
```

```
RMSE on test data: 57366.910692045196
```

```
pip install xgboost
```

```
Requirement already satisfied: xgboost in
/usr/local/lib/python3.10/dist-packages (2.0.3)
```

```
Requirement already satisfied: numpy in
/usr/local/lib/python3.10/dist-packages (from xgboost) (1.25.2)
```

```
Requirement already satisfied: scipy in
/usr/local/lib/python3.10/dist-packages (from xgboost) (1.11.4)
```

Extreme Gradient Boosting (XGBoost) is an open-source library that provides an efficient and effective implementation of the gradient boosting algorithm.

Use the scikit-learn wrapper classes: XGBRegressor and XGBClassifier.

try another machine learning algorithm : XGBoost
from xgboost import XGBRegressor

xgb_model = XGBRegressor()

Train the model using the training sets

xgb_model.fit(X_train, y_train)

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
              early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
              feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None,
              max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)
```

#run the predictions on the training and testing data

y_xgb_pred_test = xgb_model.predict(X_test)

#compare the actual values (ie, target) with the values predicted by the model

xgb_pred_test_df = pd.DataFrame({'Actual': y_test, 'Predicted':
y_xgb_pred_test})

xgb_pred_test_df

```
{"summary": "{\n  \"name\": \"xgb_pred_test_df\",\n  \"rows\": 6192,\n  \"fields\": [\n    {\n      \"column\": \"Actual\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 114575.39507173728,\n        \"min\": 14999.0,\n        \"max\": 500001.0,\n        \"num_unique_values\": 2689,\n        \"samples\": [\n          203300.0,\n          202200.0,\n          271100.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Predicted\",\n      \"properties\": {\n        \"dtype\": \"float32\",\n        \"num_unique_values\": 6189,\n        \"samples\": [\n          107010.84375,\n          119321.6640625,\n          67871.625\n        ]\n      }\n    }\n  ]\n}
```


```

],\n          \"semantic_type\": \"\", \n          \"description\": \"\"\n}\n }\n ]\n n}","type":"dataframe","variable_name":"xgb_pred_test_df"}

fig= plt.figure(figsize=(8,8))
xgb_pred_test_df = xgb_pred_test_df.reset_index()
xgb_pred_test_df = xgb_pred_test_df.drop(['index'],axis=1)
plt.plot(xgb_pred_test_df[:50])
plt.legend(['Actual value','Predicted value'])

<matplotlib.legend.Legend at 0x78b3029be8c0>

```



```

from sklearn.metrics import r2_score

```

```

score = r2_score(y_test, y_xgb_pred_test)

print("R^2 - {}".format(round(score, 2) * 100))

R^2 - 78.0%

# Determine mean square error and root mean square error
from sklearn.metrics import mean_squared_error
import math

mse = mean_squared_error(y_test, y_xgb_pred_test)
rmse = math.sqrt(mean_squared_error(y_test, y_xgb_pred_test))

print(mse)
print(rmse)

2939759040.9080276
54219.5448238735

# Calculate mean absolute error(any large error)
from sklearn.metrics import mean_absolute_error

print(mean_absolute_error(y_test, y_xgb_pred_test))

36285.050324826894

# We can build and score a model on multiple folds using cross-validation
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import cross_val_score

# define model evaluation method
cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

scores = cross_val_score(xgb_model, X, y, scoring='r2',
error_score='raise', cv=cv, n_jobs=-1, verbose=1)

#average of all the r2 scores across runs
print(scores.mean())

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.

0.7850403811484551

[Parallel(n_jobs=-1)]: Done 30 out of 30 | elapsed: 7.0s finished

# determine hyperparameter available for tuning
xgb_model.get_params()

```



```
{'objective': 'reg:squarederror',
 'base_score': None,
 'booster': None,
 'callbacks': None,
 'colsample_bylevel': None,
 'colsample_bynode': None,
 'colsample_bytree': None,
 'device': None,
 'early_stopping_rounds': None,
 'enable_categorical': False,
 'eval_metric': None,
 'feature_types': None,
 'gamma': None,
 'grow_policy': None,
 'importance_type': None,
 'interaction_constraints': None,
 'learning_rate': None,
 'max_bin': None,
 'max_cat_threshold': None,
 'max_cat_to_onehot': None,
 'max_delta_step': None,
 'max_depth': None,
 'max_leaves': None,
 'min_child_weight': None,
 'missing': nan,
 'monotone_constraints': None,
 'multi_strategy': None,
 'n_estimators': None,
 'n_jobs': None,
 'num_parallel_tree': None,
 'random_state': None,
 'reg_alpha': None,
 'reg_lambda': None,
 'sampling_method': None,
 'scale_pos_weight': None,
 'subsample': None,
 'tree_method': None,
 'validate_parameters': None,
 'verbosity': None}
```

```
xgb_model_2 = XGBRegressor(
    gamma=0.05,
    learning_rate=0.01,
    max_depth=6,
    n_estimators=1000,
    n_jobs=16,
    objective='reg:squarederror',
    subsample=0.8,
    scale_pos_weight=0,
    reg_alpha=0,
```

```

    reg_lambda=1,
    verbosity=1)

xgb_model_2.fit(X_train, y_train)

#run the predictions on the training and testing data
y_xgb_2_pred_test = xgb_model_2.predict(X_test)

# compare the actual values (ie, target) with the values predicted by the model
xgb_2_pred_test_df = pd.DataFrame({'Actual': y_test, 'Predicted':
y_xgb_2_pred_test})


xgb_2_pred_test_df

{"summary":{"\n  \"name\": \"xgb_2_pred_test_df\", \n  \"rows\": 6192, \n
  \"fields\": [\n    {\n      \"column\": \"Actual\", \n
    \"properties\": {\n      \"dtype\": \"number\", \n      \"std\":
114575.39507173728, \n      \"min\": 14999.0, \n      \"max\":
500001.0, \n      \"num_unique_values\": 2689, \n      \"samples\":
[\n        203300.0, \n        202200.0, \n        271100.0 \n
], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n
} \n    }, \n    {\n      \"column\": \"Predicted\", \n
    \"properties\": {\n      \"dtype\": \"float32\", \n
    \"num_unique_values\": 6190, \n      \"samples\": [\n
99311.6015625, \n      103087.6640625, \n      345059.84375 \n
], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n
} \n    } \n  ] \n
n} \", \"type\": \"dataframe\", \"variable_name\": \"xgb_2_pred_test_df\"}

fig= plt.figure(figsize=(8,8))
xgb_2_pred_test_df = xgb_2_pred_test_df.reset_index()
xgb_2_pred_test_df = xgb_2_pred_test_df.drop(['index'],axis=1)
plt.plot(xgb_2_pred_test_df[:50])
plt.legend(['Actual value', 'Predicted value'])

<matplotlib.legend.Legend at 0x78b3029bf0d0>

```



```
from sklearn.metrics import mean_squared_error

mse = np.sqrt(mean_squared_error(y_test, y_xgb_2_pred_test))
print("RMSE: %.2f" % (mse*(1/2.0)))

RMSE: 230.63

# Determine accuracy using R^2
r2_xgb_model_2_test = round(xgb_model_2.score(X_test, y_test), 2)

print("R^2 Test: {}".format(r2_xgb_model_2_test))

R^2 Test: 0.78
```